

# Random Forests in Practice

Data Science Lab

2025-11-21

## Table of contents

Overview . . . . .	1
Mathematical Model . . . . .	1
Environment Setup . . . . .	2
Data Loading and Preparation . . . . .	3
Model Training . . . . .	4
Diagnostics . . . . .	5
Receiver Operating Characteristic . . . . .	5
Confusion Matrix . . . . .	6
Feature Importance Visualization . . . . .	7
Hyperparameter Considerations . . . . .	8
Practical Tips . . . . .	8
References . . . . .	9

## Overview

Random forests are ensemble models that aggregate many decision trees to reduce variance and improve generalization.<sup>1</sup> This document walks through training and interpreting a random forest classifier in Python, with a mix of narrative, math, and visuals.

## Mathematical Model

Each tree  $T_b$  is trained on a bootstrap sample  $\mathcal{D}_b$  and a random subset of features. The forest prediction for a classification task with  $B$  trees is the majority vote:

$$\hat{y} = \text{mode}(\{T_b(\mathbf{x})\}_{b=1}^B)$$

---

<sup>1</sup>Introduced the random forest algorithm with theoretical justification and empirical benchmarks.

For regression, the trees are averaged:

$$\hat{y} = \frac{1}{B} \sum_{b=1}^B T_b(\mathbf{x})$$

The randomization across bootstrapped data and feature subsampling drives decorrelation between trees, delivering lower variance than single-tree models.

## Environment Setup

```
import importlib
import subprocess
import sys

def ensure(package):
    try:
        importlib.import_module(package)
    except ImportError:
        subprocess.check_call([sys.executable, "-m", "pip", "install", package])

for pkg in ("numpy", "pandas", "seaborn", "matplotlib", "scikit-learn"):
    ensure(pkg)
```

```
Requirement already satisfied: scikit-learn in /Users/luciusjmorningstar/Desktop/GIT-REPOSITORY/JJB_Gallery/.venv/lib/python3.13/site-packages (1.7.2)
Requirement already satisfied: numpy>=1.22.0 in /Users/luciusjmorningstar/Desktop/GIT-REPOSITORY/JJB_Gallery/.venv/lib/python3.13/site-packages (from scikit-learn) (2.3.5)
Requirement already satisfied: scipy>=1.8.0 in /Users/luciusjmorningstar/Desktop/GIT-REPOSITORY/JJB_Gallery/.venv/lib/python3.13/site-packages (from scikit-learn) (1.16.3)
Requirement already satisfied: joblib>=1.2.0 in /Users/luciusjmorningstar/Desktop/GIT-REPOSITORY/JJB_Gallery/.venv/lib/python3.13/site-packages (from scikit-learn) (1.5.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /Users/luciusjmorningstar/Desktop/GIT-REPOSITORY/JJB_Gallery/.venv/lib/python3.13/site-packages (from scikit-learn) (3.6.0)
```

```

import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import RocCurveDisplay, ConfusionMatrixDisplay, classification_report

sns.set_theme(style="whitegrid")

```

## Data Loading and Preparation

We will use the Breast Cancer Wisconsin dataset bundled with scikit-learn, which contains 30 features computed from digitized fine needle aspirate images.<sup>2</sup>

```

dataset = load_breast_cancer(as_frame=True)
df = dataset.frame
df.head()

```

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.30
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.086
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.197
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.243
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.198

Split the data into training and testing sets (stratified to maintain label balance).

```

X_train, X_test, y_train, y_test = train_test_split(
    df.drop(columns="target"),
    df["target"],
    test_size=0.25,
    random_state=42,
    stratify=df["target"]
)

X_train.shape, X_test.shape

```

---

<sup>2</sup>Official description of the dataset, feature definitions, and usage considerations.

```
((426, 30), (143, 30))
```

## Model Training

```
rf = RandomForestClassifier(  
    n_estimators=400,  
    max_features="sqrt",  
    min_samples_leaf=2,  
    random_state=42,  
    n_jobs=-1  
)  
rf.fit(X_train, y_train)
```

n_estimators	400
criterion	'gini'
max_depth	None
min_samples_split	2
min_samples_leaf	2
min_weight_fraction_leaf	0.0
max_features	'sqrt'
max_leaf_nodes	None
min_impurity_decrease	0.0
bootstrap	True
oob_score	False
n_jobs	-1
random_state	42
verbose	0
warm_start	False
class_weight	None
ccp_alpha	0.0
max_samples	None
monotonic_cst	None

Evaluate cross-validated training performance to estimate generalization ability.

```
cv_scores = cross_val_score(rf, X_train, y_train, cv=5)  
cv_scores.mean(), cv_scores.std()
```

```
(np.float64(0.9600547195622434), np.float64(0.020556647327639923))
```


## Diagnostics

```
y_pred = rf.predict(X_test)
print(classification_report(y_test, y_pred, target_names=dataset.target_names))
```

	precision	recall	f1-score	support
malignant	0.96	0.92	0.94	53
benign	0.96	0.98	0.97	90
accuracy			0.96	143
macro avg	0.96	0.95	0.95	143
weighted avg	0.96	0.96	0.96	143


## Receiver Operating Characteristic

```
fig, ax = plt.subplots(figsize=(6, 4))
RocCurveDisplay.from_estimator(rf, X_test, y_test, ax=ax)
ax.set_title("Random Forest ROC Curve")
plt.tight_layout()
```



### Confusion Matrix

```
fig, ax = plt.subplots(figsize=(4, 4))
ConfusionMatrixDisplay.from_estimator(rf, X_test, y_test, display_labels=dataset.target_names)
ax.set_title("Confusion Matrix")
plt.tight_layout()
```



## Feature Importance Visualization


```

importances = pd.Series(rf.feature_importances_, index=df.columns[:-1]).sort_values(ascending=False)
top_features = importances.head(15)

try:
    sns
except NameError:
    import seaborn as sns
    sns.set_theme(style="whitegrid")

plt.figure(figsize=(8, 5))
sns.barplot(x=top_features.values, y=top_features.index, palette="viridis")
plt.title("Top 15 Feature Importances")
plt.xlabel("Gini Importance")
plt.ylabel("Feature")
plt.tight_layout()

```



## Hyperparameter Considerations

- `n_estimators`: Increasing trees generally improves stability until diminishing returns set in.
- `max_depth` or `min_samples_leaf`: Control tree complexity, mitigating overfitting.
- `max_features`: Governs the degree of feature randomness; `sqrt` is typical for classification.
- `class_weight`: Useful for imbalanced datasets to penalize misclassification of minority classes.

Grid search or Bayesian optimization can systematically explore these settings.<sup>3</sup>

## Practical Tips

- **Feature scaling:** Not required because trees are invariant to monotonic transformations.
- **Missing values:** scikit-learn's implementation does not handle NaNs; impute beforehand.
- **Interpretability:** Use SHAP values or permutation importance for richer explanations.
- **Out-of-bag (OOB) estimates:** Enable `oob_score=True` to get a built-in validation metric without a separate hold-out set.

---

<sup>3</sup>Demonstrated the efficiency gains of random search over grid search for hyperparameter tuning.

## References

- Breiman, L. (2001). Random forests. *Machine Learning*, 45(1), 5–32. [https://doi.org/10.1023/A:1010933404324<sup>4</sup>](https://doi.org/10.1023/A:1010933404324)
- scikit-learn Breast Cancer Dataset docs. [https://scikit-learn.org/stable/datasets/toy\\_dataset.html<sup>5</sup>](https://scikit-learn.org/stable/datasets/toy_dataset.html)
- Bergstra, J., & Bengio, Y. (2012). Random search for hyper-parameter optimization. *Journal of Machine Learning Research*, 13, 281–305. [https://jmlr.org/papers/v13/bergstra12a.html<sup>6</sup>](https://jmlr.org/papers/v13/bergstra12a.html)

---

<sup>4</sup>breiman2001

<sup>5</sup>sklearn\_breast

<sup>6</sup>bergstra2012